

5112

11-14

Hash Tables

What is a hash table?

A dictionary that uses a hash function to determine where to store items.

a data structure supporting

Insert


Query

(Delete)

(Scan)

← not supported by
hash tables

Hash tables generalize direct mapping



In for example a DFS, can use a direct mapping to tell which nodes have been visited.

1	2	3	4	5	6	7	8	9
		x	x	x	x			

What is a hash function?


Have a universe U of keys.

e.g. integers between 0 and 2^{32}
strings of byte
points in \mathbb{R}^3 .

A hash function is a randomized function
 $h: U \longrightarrow \{0, \dots, m-1\}$
for some m .


Chaining Hash Table

Hash each item to a bucket and store each bucket as a linked list.



$$h(A) = 11 \quad h(B) = 3 \quad h(C) = 7 \quad h(D) = 3$$

Hash each item to a bucket and store each bucket as a linked list.



Could I just use random numbers?

Suppose I insert B, then query B.


No! Would the same value each time.

What properties do we want from our hash function?

Could I use a deterministic function?

Interesting question.
This is hard to analyze.

Hash each item to a bucket and store each bucket as a linked list.



What properties do we want from our hash function?

Expected cost of a query is proportional to the expected list length.

How can the hash function minimize this?
By distributing items to different buckets.

Totally Random Hash Functions

Choose h uniformly from the set of all functions from $U \rightarrow \{0, \dots, m-1\}$.

Equivalent to picking a random number from $\{0, \dots, m-1\}$ for each $x \in U$.

way too big

To store h , how many bits do I need?
log m bits for each $x \in U$, so that $|U| \log m$ bits.
Also need some way to efficiently compute.

Universal Hash Functions

Idea: pick h from a smaller set of potential hash functions.

A family of hash functions is called universal if for all $x \neq y \in \mathcal{U}$,

$$\Pr\{h(x) = h(y)\} \leq \frac{1}{m}$$

Universal Hash Functions

Example: $h_{a,b}(x) = [(ax + b) \bmod p] \bmod m.$

p is a fixed prime with $p > |U|$,
 $0 \leq a, b < p$, with $a \neq 0$.

Number theory \Rightarrow universal.

How to encode which hash function we're using:

Just need to provide p, a, b .

Universal Hash Functions

Example: Multiply-Shift.

Assume $m = 2^k$, a odd number $0 < a < 2^w$

$$h_a(x) = (ax \bmod 2^w) / 2^{w-k}$$

Multiplying x by a , then truncating to a word
Then shift right to get a k -bit result.
Not universal, but almost universal

k-wise independent hash functions

A family is k-wise independent if for all $x_1, \dots, x_k \in \mathcal{U}$,
 $t_1, \dots, t_k \in \mathcal{T}_1, \dots, \mathcal{T}_k$

$$\Pr_h \{h(x_1) = t_1 \wedge h(x_2) = t_2 \wedge \dots \wedge h(x_k) = t_k\} = O\left(\frac{1}{m^k}\right)$$

Stronger condition than universal.
Lots of constructions

Hash Functions in Practice

Basically none of the theory matters.

Murmur Hash: core loop that does a multiply and a rotation.


← determined by
a seed.

Note: These hash functions are not cryptographic.

Crypto hash functions are hard to invert.

" " " have extremely few collisions
(have > 256 bits of output)

Hash each item to a bucket and store each bucket as a linked list.



Expected cost of a query is proportional to the expected list length.

Suppose there are n items in the hash table.

Let C_t be the number of items hashing to t .


$$E[\text{cost of a query}] = E_h[C_t]$$

$$= \sum_i \Pr[h(x_i) = t]$$

$$= O\left(\frac{n}{m}\right) \quad (h \text{ is universal})$$

$$= O(1) \text{ if } m = \Omega(n)$$

Hash each item to a bucket and store each bucket as a linked list.



What about the worst case?

All items in the same bucket
 $\Rightarrow \Theta(n)$.

Very unlikely!

If h is totally random,
and $m = \Theta(n)$, then

$$C_t = O\left(\frac{\log n}{\log \log n}\right)$$


with high probability.

$\uparrow 1 - \frac{1}{nc}$ for any c .

FKS Perfect Hashing

Static: Given all n items up front

In contrast to inserting them one by one.



How large does a hash table need
to be to have no collisions?

$$\begin{aligned} E\{\# \text{ of collisions}\} &= \sum_{i < j} \Pr[h(x_i) = h(x_j)] \\ &= \frac{1}{m} \# \text{ of distinct pairs} \\ &= \frac{1}{m} \frac{n(n-1)}{2} = \Theta\left(\frac{n^2}{m}\right) \end{aligned}$$

If we choose m large enough that this
expectation is $\approx \frac{1}{2}$, then can use Markov's inequality
to say $\Pr[\text{no collisions}] \leq \frac{1}{2}$.

Markov's Inequality

If X is a non-negative R.V., and $a > 0$, then

$$\Pr\{X \geq a\} \leq \frac{E\{X\}}{a}.$$

$$\Pr\{\# \text{ collisions} \geq 1\} \leq E\{\# \text{ collisions}\} \leq \frac{1}{2}.$$